

AEROSPIKE

IN-MEMORY + NOSQL + ACID

AEROSPIKE DEPLOYMENT TESTING

Aerospike aer . o . spike [air-oh- spahyk]
noun, 1. tip of a rocket that enhances speed and stability

Objectives

At the end of this module, you will be able to:

- Load test data into an Aerospike database using the Aerospike Java Benchmark tool.
- Run benchmark tests

These tasks are ones that you can use in your environments. Doing proper benchmarking will not only help to determine what your performance is, but also useful for finding basic connectivity issues.

Benchmarking Dos And Don'ts

Do:

- Understand what you are benchmarking. Start with the most simple use case (single server, RAM storage, small objects), to make sure you understand the upper limits of the system. Then change things one at a time to bring it closer to your production configuration.
- Be aware of what you are bottlenecked on.
- Be prepared to have more power on the clients than server.
- Know if you will be impacted by network congestion.

Don't:

- Expect high performance with VMs. You must be very careful when testing in a public cloud such as Amazon.
- Expect realistic number when the client and server are on the same host.

One of the most common problems is finding network congestion.

Often this will requires network administrators to help determine if there are network issues causing the problem.

Use standard Linux commands like top to detect your bottlenecks.

The Java Benchmark Tool

- + Multi-platform
- + Portable
- + Easy to use
- + Total throughput easily seen from database nodes
- Poor performance with objects larger than 4 KB
- Must use many processes/servers to get to high throughput

<EROSPIKE

© 2014 Aerospike. All rights reserved. Confidential

4

The Java benchmark tool currently generates random text using ineffective methods. So large objects are not created quickly. Be aware of this if you want to test large objects and you will quickly saturate the CPUs on the clients.

If you want to test with large objects and need performance, you may want to use the benchmarking tool built into the C client. This is specific to Linux.

Java Benchmark Syntax

> `./run_benchmarks [options]`

Flag	Description
<code>-b, --bins <arg></code>	The number of bins (like columns) of data.
<code>-g, --throughput <arg></code>	The target total throughput. This will include all operations.
<code>-h, --host <arg></code>	A seed node (any node in the cluster). The tool will learn about the other nodes from this one.
<code>-p, --port <arg></code>	The port on the Aerospike node to connect to (default 3000)
<code>-k, --keys <arg></code>	This is the number of keys/records to use. If you select a number too small, you will get write contention. Too large and you may run out of memory.
<code>-latency <arg></code>	Produce a table of latencies as measured from the client. This can be compared with the latencies on the server to see where any bottleneck may lie. Recommended value "-latency 7,1"
<code>-n, --namespace <arg></code>	The namespace to test against. This is similar to a "database" in a relational database.
<code>-o, --objectSpec <arg></code>	The type of object to use: "I" is an integer, "S:<size>" is the size of the object. Is is recommended to start with small objects to remove the network
<code>-w, --workload <arg></code>	Type of workload: I - Insert only, test stops when the final record is loaded (based on -k above) RU, <percent> - The read percentage. "90" means a 90% read rate.
<code>-z, --threads <arg></code>	The number of threads used by the benchmark. For bare metal, 64 is a good value. For VMs, you should use 4-8 threads/core.

The Java Benchmark Tool

The Aerospike benchmark tool can be found within the full Aerospike Java Client. For this class, the tool has been pre-compiled and loaded onto the Aerospike Training AMI.

To see it, go to the appropriate directory

```
> cd /home/aerotrainig/packages/aerospike-client-java-3.0.26/benchmarks
```



Loading Test Data

Loading Data Example

Prior to running any tests, you should load data into the database. There is a script on the training server that will do this called `test_load.sh`. The active contents of this file are:

```
./run_benchmarks -h 127.0.0.1 \ # Local DB server
                  -p 3000 \      # Local port
                  -n test \      # Namespace "test"
                  -k 100000 \    # Load 100,000 records
                  -latency "7,1" \ # Show latency numbers
                  -o S:100 \     # Use 100 byte string
                  -w I \         # Insert only
                  -z 8           # Use 8 threads
```

This will take a few minutes on your VM. When the script has loaded all the data, it will exit.

You will want to do tests on reading/writing data. However, it is important to pre-load data prior to running these tests.

Loading Data Example Output

```
[aerotraining@ip-10-231-34-97 benchmarks]$ ./test_load.sh
Benchmark: 127.0.0.1:3000, namespace: test, set: <empty>, threads: 8, workload: INITIALIZE
keys: 100000, start key: 1, key length: 10, bins: 1, throughput: unlimited, debug: false
write policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500
bin type 1: string[100]
2014-08-27 22:12:52.091 INFO Thread 1 Add node BB9ED14150B0022 127.0.0.1:3000
2014-08-27 22:12:52.137 write(count=0 tps=0 timeouts=0 errors=0)
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write    0%    0%    0%    0%    0%    0%    0%
2014-08-27 22:12:53.141 write(count=2432 tps=2432 timeouts=0 errors=0)
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write   90%  10%   9%   8%   7%   5%   1%
2014-08-27 22:12:54.170 write(count=6492 tps=4060 timeouts=0 errors=0)
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write   92%   8%   7%   6%   5%   3%   1%
2014-08-27 22:12:55.194 write(count=12705 tps=6213 timeouts=0 errors=0)
...
```

While this is running, check on the AMC.

<EROSPIKE>

© 2014 Aerospike. All rights reserved. Confidential

9

Notice here that there are latencies given for writes.

The numbers represent the that exceed the timemarks at the top.



Running Benchmarks

Running Balanced Workload

The first test to run is on a balanced workload. This will test database operations that are 50% read/50% write. This is in a file called `test_balanced.sh`. The active contents of this file are:

```
./run_benchmarks -h 127.0.0.1 \ # Local DB server
                  -p 3000 \      # Local port
                  -n test \      # Namespace "test"
                  -k 100000 \    # Load 100,000 records
                  -latency "7,1" \ # Show latency numbers
                  -o S:100 \     # Use 100 byte string
                  -w RU,50 \     # Read at 50%
                  -z 8           # Use 8 threads
```

This script will continue running until it is stopped.

<EROSPIKE>

© 2014 Aerospike. All rights reserved. Confidential

11

You can run an intensive test using a balanced workload. These numbers show reading 50% of the time and updating 50% of the time.

When interpreting these remember the following:

- An update on this test is actually a read and a write. This is due to the need to update existing data. While you can replace the data rather than update, the Java benchmark tool currently only does an update.
- Keep in mind that when writing data, you will be writing a number of copies based on the replication factor. For a single node, this will be always be 1.
- You may need to run multiple copies of these scripts to maximize performance.

Running Balanced Load Example Output

```
aerotraining@ip-10-231-34-97 benchmarks]$ ./test_balanced.sh
Benchmark: 127.0.0.1:3000, namespace: test, set: <empty>, threads: 8, workload: READ_UPDATE
read: 50% (all bins: 100%, single bin: 0%), write: 50% (all bins: 100%, single bin: 0%)
keys: 100000, start key: 0, key length: 10, bins: 1, throughput: unlimited, debug: false
read policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500, reportNotFound: false
write policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500
bin type 1: string[100]
2014-08-27 22:14:56.269 INFO Thread 1 Add node BB9ED14150B0022 127.0.0.1:3000
2014-08-27 22:14:56.329 write(tps=2 timeouts=0 errors=0) read(tps=0 timeouts=0 errors=0)
total(tps=2 timeouts=0 errors=0)
<=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write 89% 11% 11% 11% 11% 5% 2%
read 86% 14% 13% 13% 13% 9% 2%
2014-08-27 22:14:57.577 write(tps=1458 timeouts=0 errors=0) read(tps=1402 timeouts=0
errors=0) total(tps=2860 timeouts=0 errors=0)
<=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write 86% 14% 13% 12% 11% 6% 1%
read 89% 11% 10% 9% 8% 5% 2%
2014-08-27 22:14:58.585 write(tps=1994 timeouts=0 errors=0) read(tps=1934 timeouts=0
errors=0) total(tps=3928 timeouts=0 errors=0)
<=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write 93% 7% 7% 6% 5% 3% 1%
read 94% 6% 5% 5% 4% 2% 1%
...
```

While this is running, check on the AMC.

Running Balanced Workload

You can then run a workload based on high read rates. This will test database operations that are 95% read/5% write. This is in a file called `test_highread.sh`. The active contents of this file are:

```
./run_benchmarks -h 127.0.0.1 \ # Local DB server
                  -p 3000 \      # Local port
                  -n test \      # Namespace "test"
                  -k 100000 \    # Load 100,000 records
                  -latency "7,1" \ # Show latency numbers
                  -o S:100 \     # Use 100 byte string
                  -w RU,95 \     # Read at 95%
                  -z 8           # Use 8 threads
```

This script will continue running until it is stopped.

Running Balanced Load Example Output

```
[aerotraining@ip-10-231-34-97 benchmarks]$ ./test_highread.sh
Benchmark: 127.0.0.1:3000, namespace: test, set: <empty>, threads: 8, workload: READ_UPDATE
read: 95% (all bins: 100%, single bin: 0%), write: 5% (all bins: 100%, single bin: 0%)
keys: 100000, start key: 0, key length: 10, bins: 1, throughput: unlimited, debug: false
read policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500, reportNotFound: false
write policy: timeout: 0, maxRetries: 2, sleepBetweenRetries: 500
bin type 1: string[100]
2014-08-27 22:18:08.031 INFO Thread 1 Add node BB9ED14150B0022 127.0.0.1:3000
2014-08-27 22:18:08.121 write(tps=7 timeouts=0 errors=0) read(tps=68 timeouts=0 errors=0)
total(tps=75 timeouts=0 errors=0)
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write  57%  43%  43%  29%  14%   0%   0%
read   91%   9%   9%   9%   7%   3%   0%
2014-08-27 22:18:09.130 write(tps=107 timeouts=0 errors=0) read(tps=2650 timeouts=0
errors=0) total(tps=2757 timeouts=0 errors=0)
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write  85%  15%  14%  11%  11%   4%   2%
read   90%  10%   9%   8%   7%   4%   1%
2014-08-27 22:18:10.179 write(tps=263 timeouts=0 errors=0) read(tps=5012 timeouts=0
errors=0) total(tps=5275 timeouts=0 errors=0)
    <=1ms >1ms >2ms >4ms >8ms >16ms >32ms
write  93%   7%   5%   5%   5%   3%   2%
read   94%   6%   5%   5%   4%   2%   1%
...
```

While this is running, check on the AMC.

Summary

What we have covered

- Load test data into an Aerospike database using the Aerospike Java Benchmark tool.
- Run benchmark tests

You can use the benchmark tool by simply copying the files and running them from different servers.